

```

> library(faraway)
> library(psych)
> library(caret)
> library(tidyverse)
> library(pROC)
> library(class)
> library(e1071)
> library(reprex)
> library(DMwR)
> library(vegan)
> #Data Preprocessing
> str(pima)
'data.frame': 768 obs. of 9 variables:
 $ pregnant : int 6 1 8 1 0 5 3 10 2 8 ...
 $ glucose : int 148 85 183 89 137 116 78 115 197 125 ...
 $ diastolic: int 72 66 64 66 40 74 50 0 70 96 ...
 $ triceps : int 35 29 0 23 35 0 32 0 45 0 ...
 $ insulin : int 0 0 0 94 168 0 88 0 543 0 ...
 $ bmi : num 33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
 $ diabetes : num 0.627 0.351 0.672 0.167 2.288 ...
 $ age : int 50 31 32 21 33 30 26 29 53 54 ...
 $ test : int 1 0 1 0 1 0 1 0 1 1 ...
> pima$glucose[pima$glucose == 0] = NA
> pima$diastolic[pima$diastolic == 0] = NA
> pima$triceps[pima$triceps == 0] = NA
> pima$insulin[pima$insulin == 0] = NA
> pima$bmi[pima$bmi == 0] = NA
> pima$diabetes[pima$diabetes == 0] = NA
> pima$age[pima$age == 0] = NA
> sapply(pima, function(x) sum(is.na(x)))
pregnant glucose diastolic triceps insulin bmi diabetes
age test
0 0 5 35 227 374 11 0
0 0
> pima = pima[-manyNAs(pima),]
> sapply(pima, function(x) sum(is.na(x)))
pregnant glucose diastolic triceps insulin bmi diabetes
age test
0 0 1 0 0 140 1 0
0 0
> pima.clean = knnImputation(pima, k=10)
> pima.clean$test = as.factor(pima.clean$test)
> levels(pima.clean$test) = c("No", "Yes")
> table(pima.clean$test)

```

```

No Yes
357 177
> #Creating the testing and training dataset
> set.seed(123)
> pima.train = sample_frac(tbl = pima.clean, replace=FALSE,
size=0.75)
> pima.test = anti_join(pima.clean, pima.train)
Joining, by = c("pregnant", "glucose", "diastolic", "triceps",
"insulin", "bmi", "diabetes", "age", "test")
> pima.train.lab = pima.train[,9]
> pima.test.lab = pima.test[,9]
> pima.train.data = pima.train[,1:8]
> pima.test.data = pima.test[,1:8]
> #Let us run a BLR first
> BLR = glm(test~., family=binomial, data=pima.train)
> summary(BLR)

```

Call:

```
glm(formula = test ~ ., family = binomial, data = pima.train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.8811	-0.6569	-0.3811	0.6329	2.4577

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-9.0860085	1.1663236	-7.790	6.69e-15	***
pregnant	0.1371466	0.0507069	2.705	0.00684	**
glucose	0.0332173	0.0055084	6.030	1.64e-09	***
diastolic	-0.0083542	0.0118990	-0.702	0.48262	
triceps	0.0054582	0.0171733	0.318	0.75061	
insulin	0.0004141	0.0013949	0.297	0.76660	
bmi	0.0779741	0.0266807	2.922	0.00347	**
diabetes	1.1084110	0.3942004	2.812	0.00493	**
age	0.0241565	0.0167812	1.440	0.15001	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 498.43 on 399 degrees of freedom  
Residual deviance: 351.31 on 391 degrees of freedom  
AIC: 369.31

Number of Fisher Scoring iterations: 5

```

> #Improved BLR
> BLR2 = glm(test~pregnant+glucose+bmi+diabetes, family=binomial,
data=pima.train)
> summary(BLR2)

Call:
glm(formula = test ~ pregnant + glucose + bmi + diabetes, family =
binomial,
     data = pima.train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.8979  -0.6509  -0.3960   0.6130   2.5165

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -9.144281   1.005107  -9.098  < 2e-16 ***
pregnant      0.180351   0.039704   4.542 5.56e-06 ***
glucose       0.035315   0.004633   7.622 2.50e-14 ***
bmi           0.078391   0.020840   3.762 0.000169 ***
diabetes      1.190484   0.389996   3.053 0.002269 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 498.43  on 399  degrees of freedom
Residual deviance: 353.94  on 395  degrees of freedom
AIC: 363.94

Number of Fisher Scoring iterations: 5

> AIC(BLR, BLR2)
      df      AIC
BLR    9 369.3126
BLR2   5 363.9405
> #ROC Curve
> test.prob = predict(BLR2, newdata=pima.test, type="response")
> test.roc = roc(pima.test$test~test.prob, plot=TRUE, print.auc =
TRUE)
Setting levels: control = No, case = Yes
Setting direction: controls < cases
> #Normalize the data first
> pima.train.norm = decostand(pima.train.data, "normalize")

```

```
> pima.test.norm = decostand(pima.train.data, "normalize")
> #let's try using CV, then train the model
> pima.cv = trainControl(method="repeatedcv", number=10, repeats=6)
> train.knn = train(test~., data=pima.train, method="knn", trControl
= pima.cv)
> train.knn
```

k-Nearest Neighbors

```
400 samples
  8 predictor
  2 classes: 'No', 'Yes'
```

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 6 times)

Summary of sample sizes: 359, 359, 360, 361, 360, 359, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.7620672	0.4295923
7	0.7548562	0.4224442
9	0.7507312	0.4063231

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was k = 5.

```
> #optimal k value = 5 (accuracy of 76.206%)
> pima.knn = knn(pima.train.norm, pima.test.norm, cl=pima.train.lab,
k=5, prob=TRUE)
> summary(pima.knn)
  No Yes
298 102
```

```
> confusionMatrix(pima.knn, pima.train.lab)
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	252	46
Yes	22	80

```

              Accuracy : 0.83
              95% CI   : (0.7895, 0.8655)
No Information Rate : 0.685
P-Value [Acc > NIR] : 3.198e-11
```

Kappa : 0.5847

McNemar's Test P-Value : 0.005285

Sensitivity : 0.9197  
Specificity : 0.6349  
Pos Pred Value : 0.8456  
Neg Pred Value : 0.7843  
Prevalence : 0.6850  
Detection Rate : 0.6300  
Detection Prevalence : 0.7450  
Balanced Accuracy : 0.7773

'Positive' Class : No

```
> #Now let's try a SVM  
> x.pima = subset(pima.train, select=-test)  
> y.pima = pima.train.lab  
> pima.svm = svm(x.pima, y.pima, kernel="linear")  
> summary(pima.svm)
```

Call:

```
svm.default(x = x.pima, y = y.pima, kernel = "linear")
```

Parameters:

SVM-Type: C-classification  
SVM-Kernel: linear  
cost: 1

Number of Support Vectors: 192

( 97 95 )

Number of Classes: 2

Levels:

No Yes

```

> pima.svm.pred = predict(pima.svm, x.pima)
> pima.tab = table(pima.svm.pred, y.pima)
> prop.table(pima.tab)
      y.pima
pima.svm.pred  No   Yes
      No  0.625 0.140
      Yes 0.060 0.175
> mean(pima.svm.pred == y.pima)
[1] 0.8
> #Let's use a CV (tune) to optimize the model
> pima.svm.tune <- tune(svm, x.pima, y.pima, kernel = "radial",
ranges = list(cost = 1^(-1:2), gamma = c(.5,1,2)))
> summary(pima.svm.tune)

```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```

cost gamma
  1    0.5

```

- best performance: 0.2275

- Detailed performance results:

	cost	gamma	error	dispersion
1	1	0.5	0.2275	0.06061032
2	1	0.5	0.2275	0.06061032
3	1	0.5	0.2275	0.06061032
4	1	0.5	0.2275	0.06061032
5	1	1.0	0.2975	0.05706965
6	1	1.0	0.2975	0.05706965
7	1	1.0	0.2975	0.05706965
8	1	1.0	0.2975	0.05706965
9	1	2.0	0.3175	0.05779514
10	1	2.0	0.3175	0.05779514
11	1	2.0	0.3175	0.05779514
12	1	2.0	0.3175	0.05779514

```

> pima.svm.2 = svm(x.pima, y.pima, kernel="radial",cost=1, gamma=0.5)
> summary(pima.svm.2)

```

Call:

```

svm.default(x = x.pima, y = y.pima, kernel = "radial", gamma = 0.5,
cost = 1)

```

Parameters:

SVM-Type: C-classification  
SVM-Kernel: radial  
cost: 1

Number of Support Vectors: 297

( 171 126 )

Number of Classes: 2

Levels:

No Yes

```
> pima.svm.pred2 = predict(pima.svm.2, x.pima)
> pima.svm.table = table(pima.svm.pred2, y.pima)
> prop.table(pima.svm.table)
```

	y.pima	
pima.svm.pred2	No	Yes
No	0.680	0.055
Yes	0.005	0.260

```
> mean(pima.svm.pred2 == y.pima)
[1] 0.94
```